

HW2

Problem 1 (Baseline for Variance Reduction)

a. By Monte Carlo policy gradient, we can sample a trajectory τ to estimate V^{π_θ} . In this problem, there are only three possible trajectory, which are:

1. $s \rightarrow a \rightarrow r(s, a) \rightarrow s_1$ (terminal), simply denoted as a
2. $s \rightarrow b \rightarrow r(s, b) \rightarrow s_1$ (terminal), simply denoted as b
3. $s \rightarrow c \rightarrow r(s, c) \rightarrow s_1$ (terminal), simply denoted as c

Let's calculate score function for each trajectory first:

$$\begin{aligned}\nabla_\theta \log \pi_\theta(a|s) &= \nabla_\theta \theta_a - \log(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \\ &= \begin{bmatrix} 1 - e^{\theta_a}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \\ -e^{\theta_b}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \\ -e^{\theta_c}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \end{bmatrix}\end{aligned}$$

Similarly for action b and c , then we can get:

$$\begin{aligned}\nabla_\theta \log \pi_\theta(b|s) &= \begin{bmatrix} -e^{\theta_a}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \\ 1 - e^{\theta_b}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \\ -e^{\theta_c}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \end{bmatrix} \\ \nabla_\theta \log \pi_\theta(c|s) &= \begin{bmatrix} -e^{\theta_a}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \\ -e^{\theta_b}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \\ 1 - e^{\theta_c}/(e^{\theta_a} + e^{\theta_b} + e^{\theta_c}) \end{bmatrix}\end{aligned}$$

We already know that $\theta_a = 0$, $\theta_b = \ln 5$, $\theta_c = \ln 4$, so:

$$\hat{\nabla}_a = 100 \cdot \begin{bmatrix} 1 - 0.1 \\ -0.5 \\ -0.4 \end{bmatrix} = \begin{bmatrix} 90 \\ -50 \\ -40 \end{bmatrix}, \quad \hat{\nabla}_b = 98 \cdot \begin{bmatrix} -0.1 \\ 1 - 0.5 \\ -0.4 \end{bmatrix} = \begin{bmatrix} -9.8 \\ 49 \\ -39.2 \end{bmatrix}, \quad \hat{\nabla}_c = 95 \cdot \begin{bmatrix} -0.1 \\ -0.5 \\ 1 - 0.4 \end{bmatrix} = \begin{bmatrix} -9.5 \\ -47.5 \\ 57 \end{bmatrix}$$

Then mean vector of $\hat{\nabla}V$ is:

$$\mathbb{E}[\hat{\nabla}V] = 0.1 \cdot \begin{bmatrix} 90 \\ -50 \\ -40 \end{bmatrix} + 0.5 \cdot \begin{bmatrix} -9.8 \\ 49 \\ -39.2 \end{bmatrix} + 0.4 \cdot \begin{bmatrix} -9.5 \\ -47.5 \\ 57 \end{bmatrix} = \begin{bmatrix} 0.3 \\ 0.5 \\ -0.8 \end{bmatrix}$$

To compute covariance matrix of $\hat{\nabla}V$, we need to compute $\sum_\tau P(\tau)(\hat{\nabla}_\tau - \mathbb{E}[\hat{\nabla}V])(\hat{\nabla}_\tau - \mathbb{E}[\hat{\nabla}V])^T$, here I run the following python code to help me calculate the value.

```
import numpy as np
a = np.array([90, -50, -40])
b = np.array([-9.8, 49, -39.2])
c = np.array([-9.5, -47.5, 57])
E = np.array([0.3, 0.5, -0.8])
Cov = 0.1 * np.outer(a - E, a - E) + 0.5 * np.outer(b - E, b - E) + 0.4 * np.outer(c - E, c - E)
print(Cov)
```

Output:

```
[[ 894.03  -509.75  -384.28]
 [ -509.75  2352.75 -1843.  ]
 [ -384.28 -1843.    2227.28]]
```

So the covariance matrix of $\hat{\nabla}V$ is:

$$\begin{bmatrix} 894.03 & -509.75 & -384.28 \\ -509.75 & 2352.75 & -1843 \\ -384.28 & -1843 & 2227.84 \end{bmatrix}$$

b. When we introduce baseline $B(s)$ to policy gradient, it won't change the expectation. So the expectation of $\tilde{\nabla}V$ is also:

$$\mathbb{E}[\tilde{\nabla}V] = \begin{bmatrix} 0.3 \\ 0.5 \\ -0.8 \end{bmatrix}$$

Here we choose $B(s) = V^{\pi_\theta}(s) = 0.1 \cdot 100 + 0.5 \cdot 98 + 0.4 \cdot 95 = 97$, then for each trajectory:

$$\tilde{\nabla}_a = (100 - 97) \cdot \begin{bmatrix} 0.9 \\ -0.5 \\ -0.4 \end{bmatrix} = \begin{bmatrix} 2.7 \\ -1.5 \\ -1.2 \end{bmatrix}, \quad \tilde{\nabla}_b = (98 - 97) \cdot \begin{bmatrix} -0.1 \\ 0.5 \\ -0.4 \end{bmatrix} = \begin{bmatrix} -0.1 \\ 0.5 \\ -0.4 \end{bmatrix}, \quad \tilde{\nabla}_c = (95 - 97) \cdot \begin{bmatrix} -0.1 \\ -0.5 \\ 0.6 \end{bmatrix} = \begin{bmatrix} 0.2 \\ 1 \\ -1.2 \end{bmatrix}$$

Similar to Part A, I run the following python code to help me calculate the value.

```
import numpy as np
a_ = np.array([2.7, -1.5, -1.2])
b_ = np.array([-0.1, 0.5, -0.4])
c_ = np.array([0.2, 1, -1.2])
E = np.array([0.3, 0.5, -0.8])
Cov_ = 0.1 * np.outer(a_ - E, a_ - E) + 0.5 * np.outer(b_ - E, b_ - E) + 0.4 * np.outer(c_ - E, c_ - E)
print(Cov_)
```

Output:

```
[[ 0.66 -0.5 -0.16]
 [-0.5  0.5  0. ]
 [-0.16 0.  0.16]]
```

So the covariance matrix of $\tilde{\nabla}V$ is:

$$\begin{bmatrix} 0.66 & -0.5 & -0.16 \\ -0.5 & 0.5 & 0 \\ -0.16 & 0 & 0.16 \end{bmatrix}$$

c. Use python package `scipy.optimize` to find the optimal base:

```
import numpy as np
from scipy.optimize import minimize

def trace_of_cov(base):
    a = (100 - base) * np.array([0.9, -0.5, -0.4])
    b = (98 - base) * np.array([-0.1, 0.5, -0.4])
    c = (95 - base) * np.array([-0.1, -0.5, 0.6])
    return np.trace(0.1 * np.outer(a - E, a - E) + 0.5 * np.outer(b - E, b - E) + 0.4 * np.outer(c - E, c - E))

base = minimize(trace_of_cov, x0=97)
print(f"Optimal base B(s) = {round(base.x[0],3)}")
```

Output:

```
Optimal base B(s) = 97.138
```

Problem 2 (Non-Uniform Polyak-Lojacsiewicz Condition in RL)

a. **Proof:**

$$\begin{aligned} \left\| \frac{\partial V^{\pi_\theta}(\mu)}{\partial \theta} \right\|_2 &= \left[\sum_{s,a} \left(\frac{\partial V^{\pi_\theta}(\mu)}{\partial \theta(s,a)} \right)^2 \right]^{\frac{1}{2}} \\ &\geq \left[\sum_s \left(\frac{\partial V^{\pi_\theta}(\mu)}{\partial \theta(s, a^*(s))} \right)^2 \right]^{\frac{1}{2}} \\ &\geq \frac{1}{\sqrt{S}} \sum_s \left| \frac{\partial V^{\pi_\theta}(\mu)}{\partial \theta(s, a^*(s))} \right| \quad (\text{by Cauchy-Schwarz, } \|x\|_1 \leq \sqrt{S} \|x\|_2) \\ &= \frac{1}{1-\gamma} \cdot \frac{1}{\sqrt{S}} \sum_s |d_\mu^{\pi_\theta}(s) \cdot \pi_\theta(a^*(s)|s) \cdot A^{\pi_\theta}(s, a^*(s))| \quad (\text{PG under softmax policy parametrization}) \\ &= \frac{1}{1-\gamma} \cdot \frac{1}{\sqrt{S}} \sum_s d_\mu^{\pi_\theta}(s) \cdot \pi_\theta(a^*(s)|s) \cdot |A^{\pi_\theta}(s, a^*(s))| \quad (\text{because } d_\mu^{\pi_\theta}(s) \geq 0 \text{ and } \pi_\theta(a^*(s)|s) \geq 0) \end{aligned}$$

b. **Proof:**

Define the distribution mismatch coefficient as $\alpha = \max_s \frac{d_\mu^{\pi^*}(s)}{d_\mu^{\pi_\theta}(s)}$. We have,

$$\begin{aligned}
\|\nabla V^{\pi_\theta}(\mu)\|_2 &\geq \frac{1}{1-\gamma} \cdot \frac{1}{\sqrt{S}} \sum_s \frac{d_\mu^{\pi_\theta}(s)}{d_\rho^{\pi^*}(s)} \cdot d_\rho^{\pi^*}(s) \cdot \pi_\theta(a^*(s)|s) \cdot |A^{\pi_\theta}(s, a^*(s))| \\
&\geq \frac{1}{1-\gamma} \cdot \frac{1}{\sqrt{S}} \cdot \left\| \frac{d_\rho^{\pi^*}}{d_\mu^{\pi_\theta}} \right\|_\infty^{-1} \cdot \min_s \pi_\theta(a^*(s)|s) \cdot \sum_s d_\rho^{\pi^*}(s) \cdot |A^{\pi_\theta}(s, a^*(s))| \\
&\geq \frac{1}{1-\gamma} \cdot \frac{1}{\sqrt{S}} \cdot \left\| \frac{d_\rho^{\pi^*}}{d_\mu^{\pi_\theta}} \right\|_\infty^{-1} \cdot \min_s \pi_\theta(a^*(s)|s) \cdot \sum_s d_\rho^{\pi^*}(s) \cdot A^{\pi_\theta}(s, a^*(s)) \\
&= \frac{1}{\sqrt{S}} \cdot \left\| \frac{d_\rho^{\pi^*}}{d_\mu^{\pi_\theta}} \right\|_\infty^{-1} \cdot \min_s \pi_\theta(a^*(s)|s) \cdot \frac{1}{1-\gamma} \sum_s d_\rho^{\pi^*}(s) \cdot \sum_a \pi^*(a|s) \cdot A^{\pi_\theta}(s, a) \\
&= \frac{1}{\sqrt{S}} \cdot \left\| \frac{d_\rho^{\pi^*}}{d_\mu^{\pi_\theta}} \right\|_\infty^{-1} \cdot \min_s \pi_\theta(a^*(s)|s) \cdot [V^*(\rho) - V^{\pi_\theta}(\rho)]
\end{aligned}$$

Problem 3 (Monte Carlo Policy Evaluation)

Property 1: Show that the true value function at state S (denoted by $V(S)$) satisfies that:

$$V(S) = \frac{P_S}{P_T} R_S + R_T$$

Proof.

$$V(S) = (P_S + P_S P_S + P_S P_S P_S + \dots) R_S + R_T = \frac{P_S}{1 - P_S} R_S + R_T = \frac{P_S}{P_T} R_S + R_T$$

Property 2: Suppose we construct an every-visit MC estimate based on only 1 trajectory τ (denoted by $\hat{V}_{MC}(S; \tau)$). Then, please show that

$$\mathbb{E}[\hat{V}_{MC}(S; \tau)] = \frac{P_S}{2P_T} R_S + R_T$$

Proof.

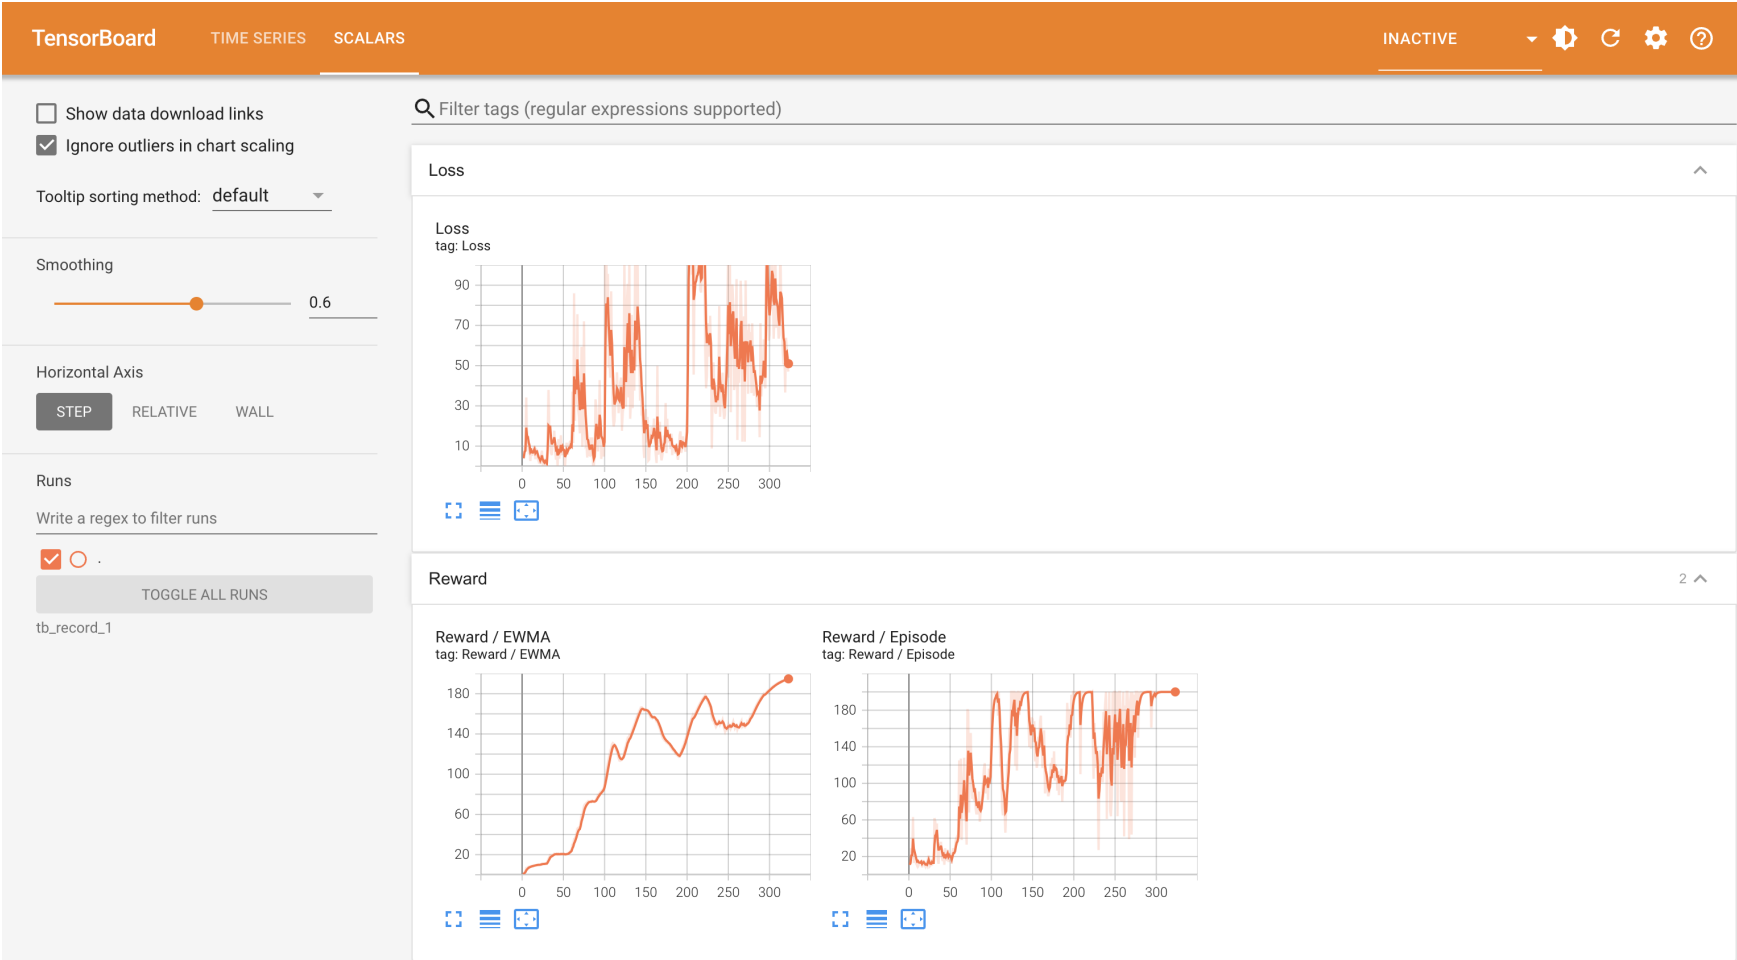
Consider all possible trajectories and the corresponding probabilities, we can get:

$$\begin{aligned}
\mathbb{E}[\hat{V}_{MC}(S; \tau)] &= \sum_{k=0}^{\infty} P_T P_S^k \left(\frac{R_S + 2R_S + \dots + kR_S + (k+1)R_T}{k+1} \right) \\
&= \sum_{k=0}^{\infty} P_T P_S^k \left(\frac{k \cdot (k+1) \cdot R_S}{2 \cdot (k+1)} + R_T \right) \\
&= \sum_{k=0}^{\infty} P_T P_S^k \left(\frac{k \cdot R_S}{2} + R_T \right) \\
&= \frac{P_T P_S R_S}{2} \sum_{k=0}^{\infty} k P_S^{k-1} + \sum_{k=0}^{\infty} P_T P_S^k R_T \\
&= \frac{P_T P_S R_S}{2(1 - P_S)^2} + P_T \cdot \frac{1}{1 - P_S} R_T \\
&= \frac{P_T P_S R_S}{2(P_T)^2} + P_T \cdot \frac{1}{P_T} R_T \\
&= \frac{P_S}{2P_T} R_S + R_T
\end{aligned}$$

Problem 4 (Policy Gradient Algorithms With Function Approximation)

Vanilla REINFORCE

- NN architecture
 - Layer1: Shared layer, `nn.Linear(self.observation_dim, 128)` then `ReLU()`
 - For Actor: `nn.Sequential(nn.Linear(128, self.action_dim), nn.Softmax(dim=-1))`
 - For Critic: `nn.Linear(128, 1)`
- Learning rate: `lr = 0.01`

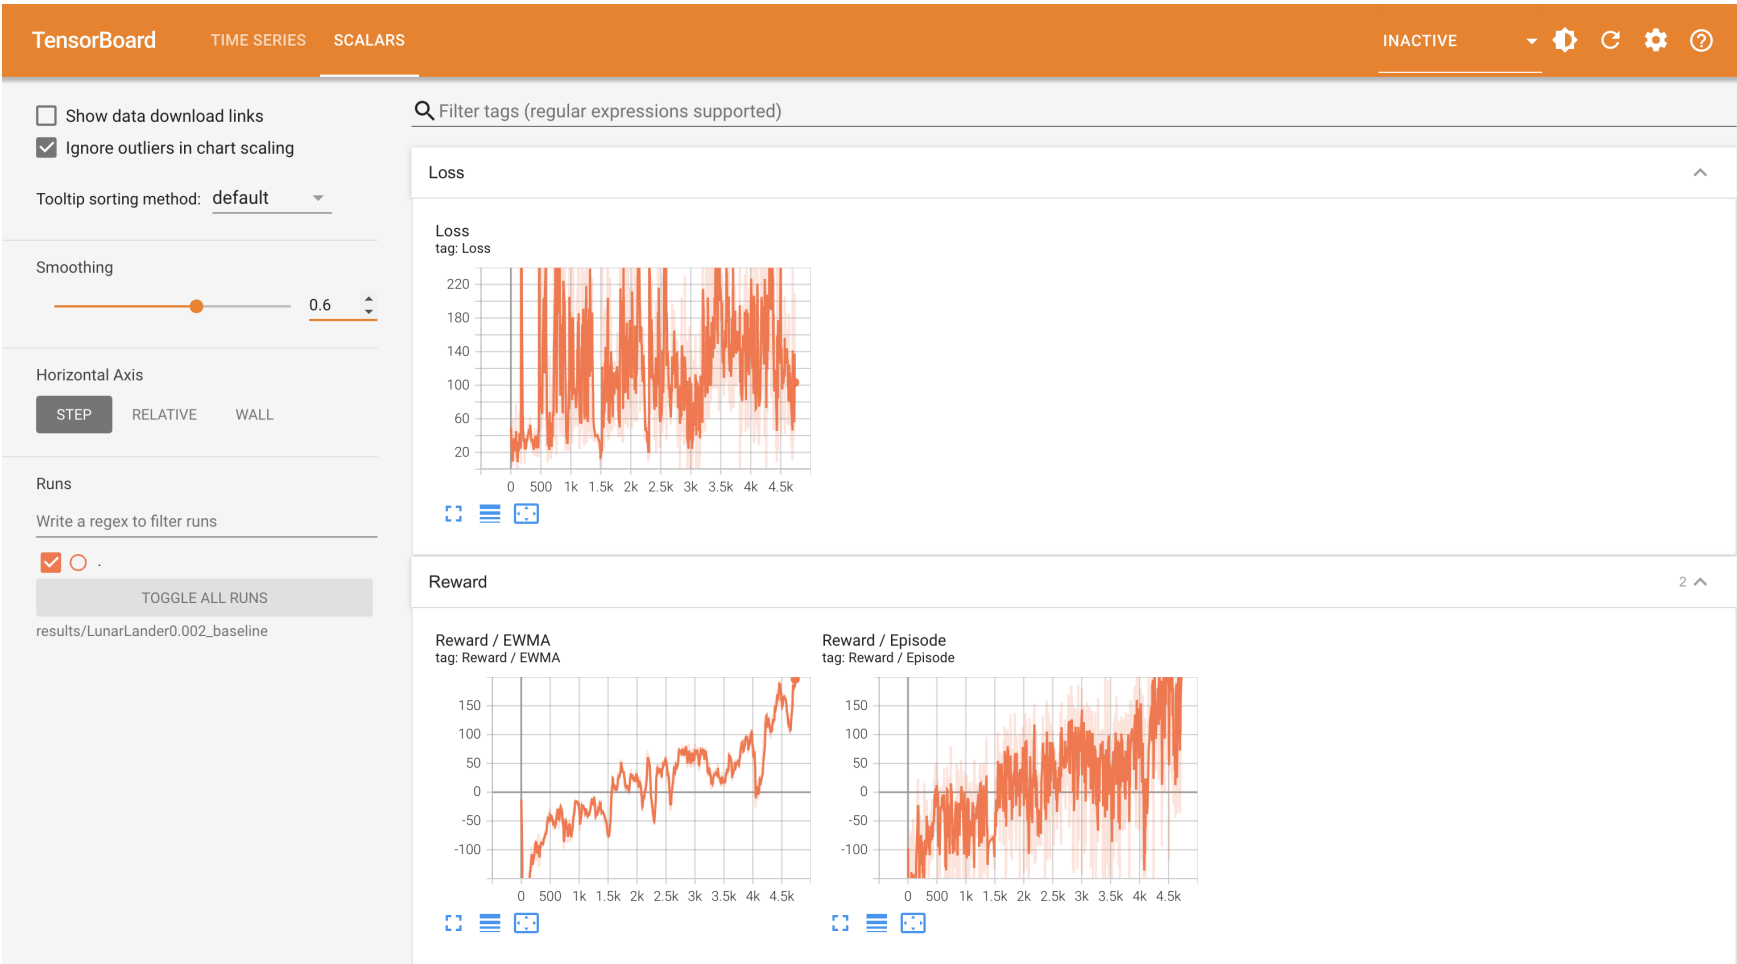


REINFORCE with baseline

- Choose state value as baseline
- NN architecture
 - Share layer: `nn.Sequential`

```
self.affine1 = nn.Sequential(
    nn.Linear(self.observation_dim, self.hidden_size),
    nn.ReLU(),
    nn.Linear(self.hidden_size, self.hidden_size),
    nn.ReLU()
)
```

- For Actor: `nn.Sequential(nn.Linear(128, self.action_dim), nn.Softmax(dim=-1))`
- For Critic: `nn.Linear(128, 1)`
- Learning rate: `lr = 0.002`



REINFORCE with GAE

- NN architecture
 - The same with NN use in *REINFORCE with baseline*
 - Share layer: `nn.Sequential`

```
self.affine1 = nn.Sequential(  
    nn.Linear(self.observation_dim, self.hidden_size),  
    nn.ReLU(),  
    nn.Linear(self.hidden_size, self.hidden_size),  
    nn.ReLU()  
)
```

- For Actor: `nn.Sequential(nn.Linear(128, self.action_dim), nn.Softmax(dim=-1))`
- For Critic: `nn.Linear(128, 1)`

• Learning rate: `lr = 0.002`

• Lambda: `0.99`

• Advantage Calculation Process

The process of calculating the advantage estimates is as follows:

1. **Initialization:** An empty list `advantages` is initialized to store the advantage values. Two variables, `advantage` and `next_value`, are initialized to zero. The `advantage` variable accumulates the discounted TD-errors, and `next_value` holds the value estimate for the state at the next time step.
2. **Backward Pass Through Time:** The function iterates backward through each time step:
 - For each step `i`, it retrieves the reward `r` and the current value estimate `v`.
 - The Temporal Difference (TD) error is calculated as:
$$\text{td_error} = r + (\gamma \times \text{next_value}) - v$$
 - The `advantage` is updated using the TD-error, discounted by both `gamma` and `lambda`:
$$\text{advantage} = \text{td_error} + (\gamma \times \lambda \times \text{advantage})$$
 - This updated advantage is prepended to the `advantages` list to maintain the correct order (since the iteration is backward).
 - The `next_value` is updated to the current value `v` for use in the next iteration.
3. **Tensor Conversion:** Finally, the list of advantage values is converted to a PyTorch tensor for compatibility with other calculations in neural networks or for use with gradient-based optimization algorithms.

